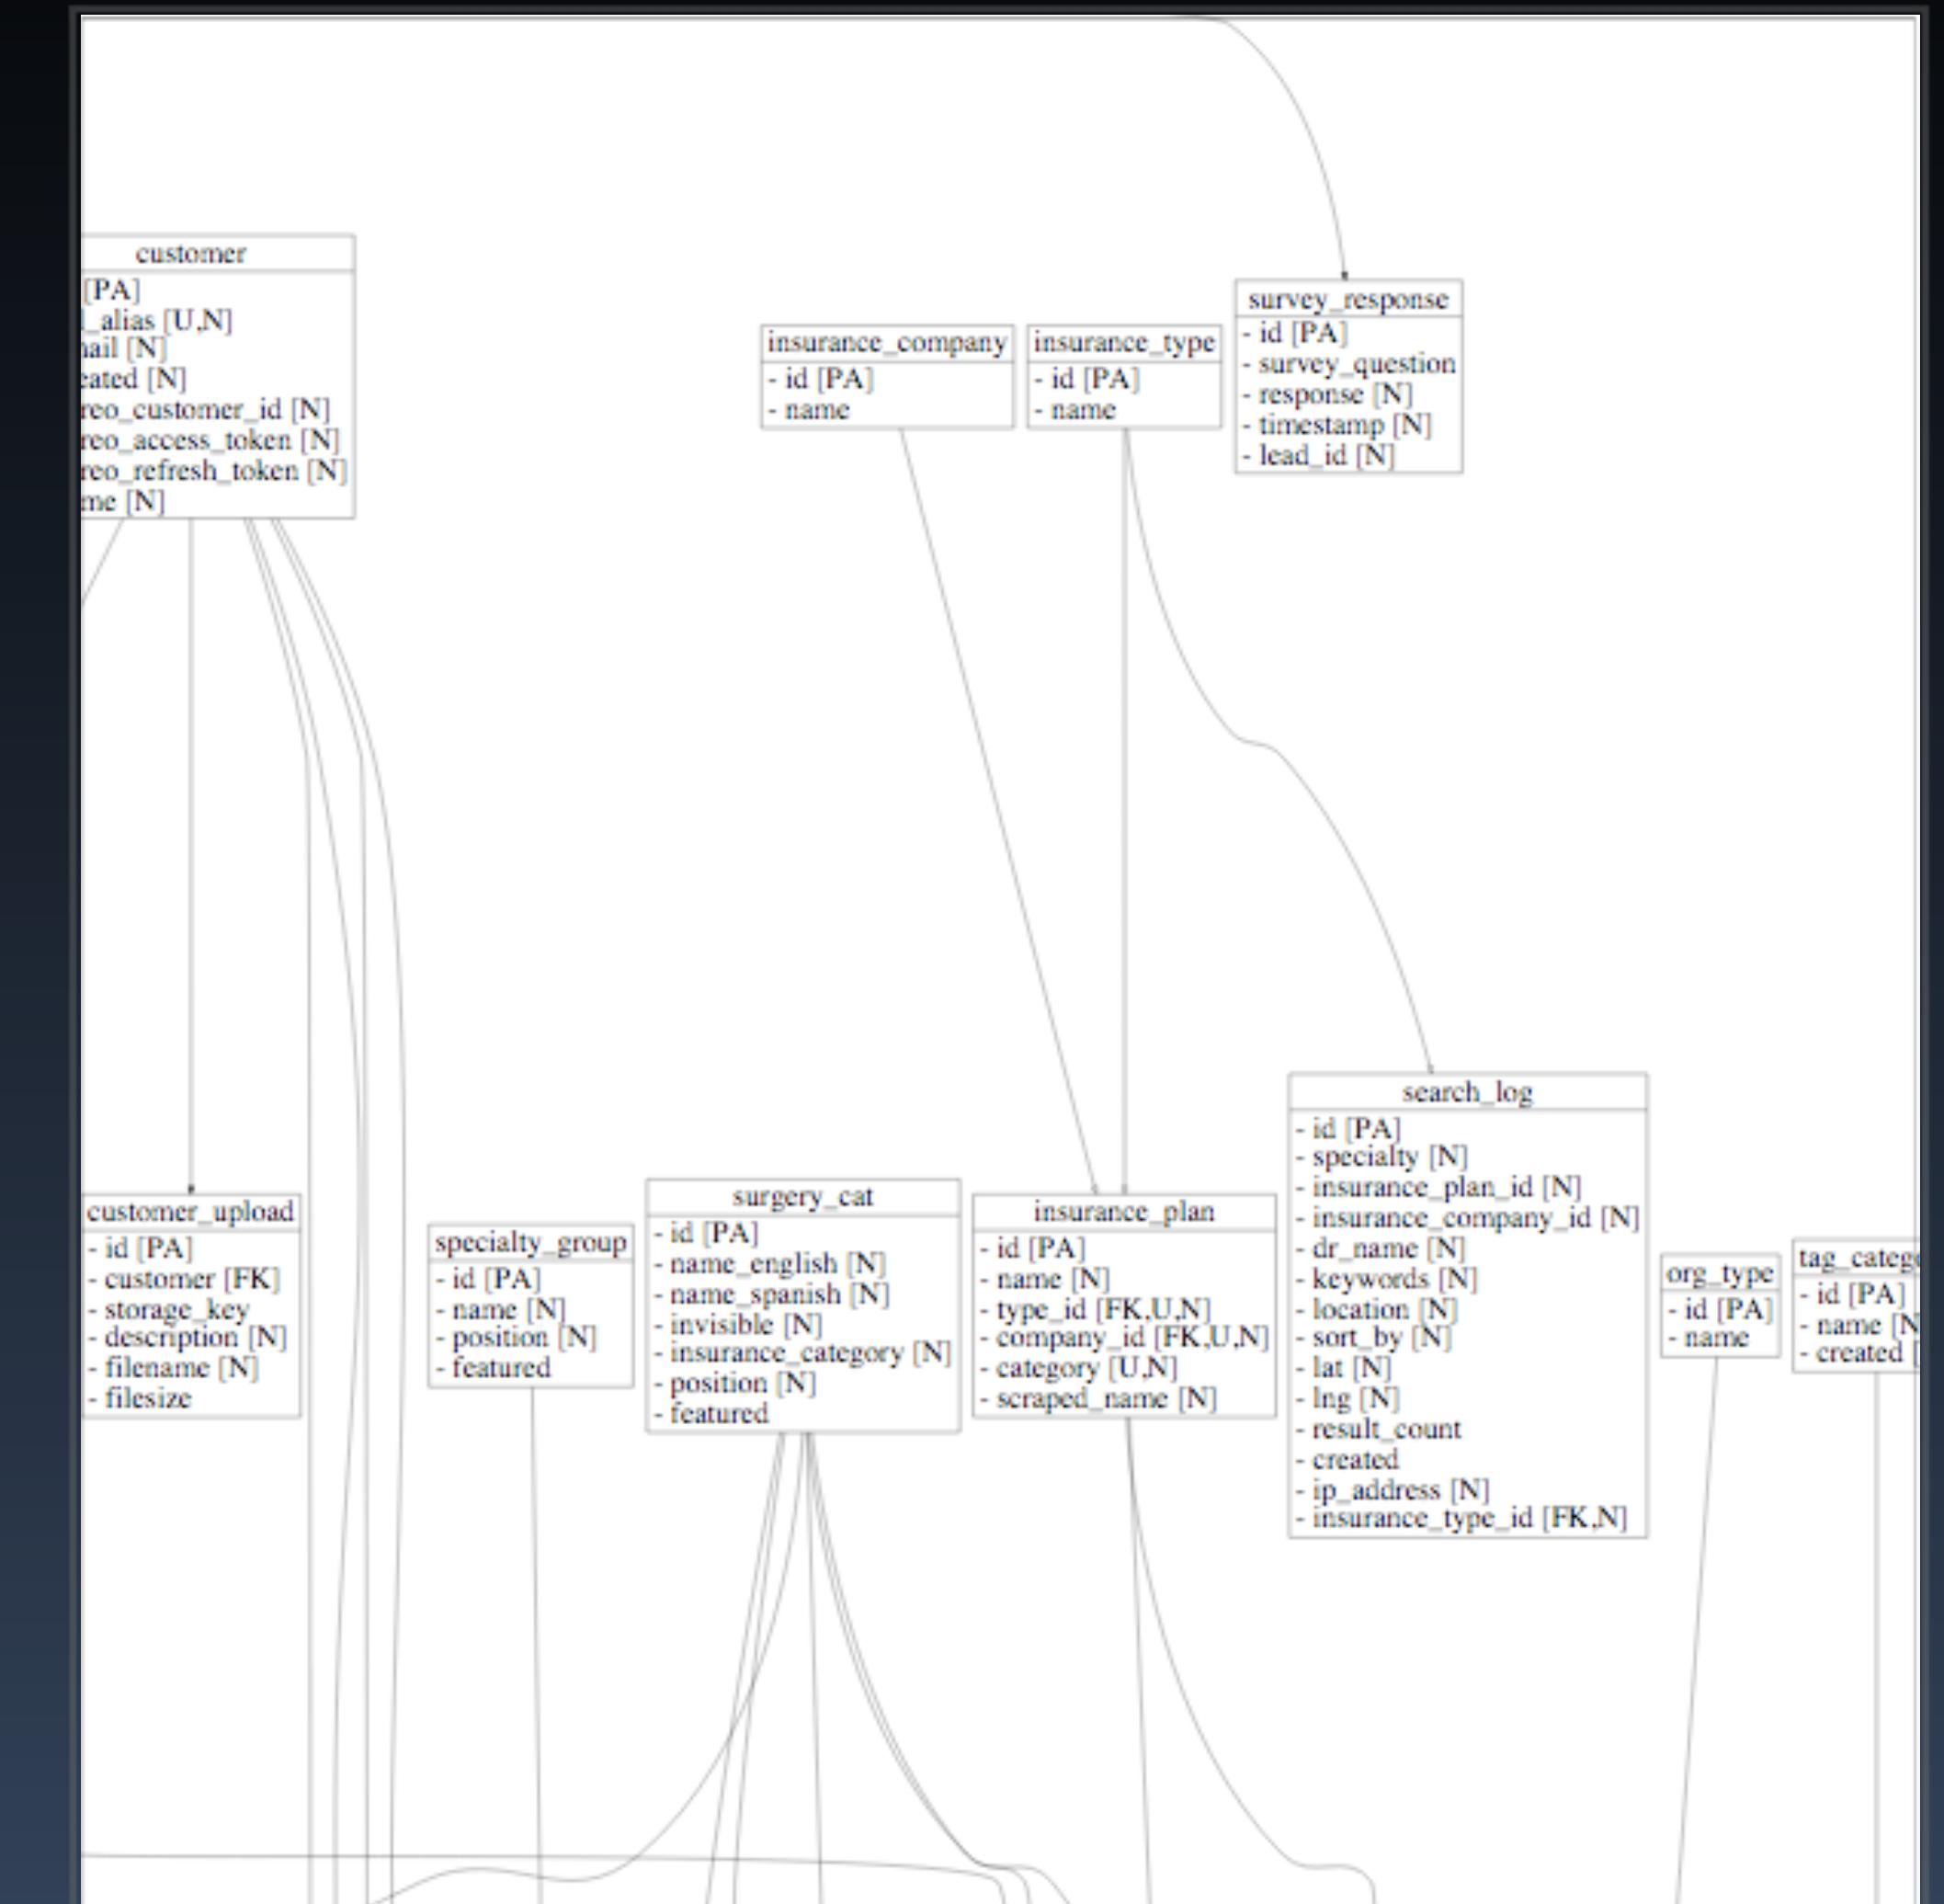


SCHEMA DESIGN MISTAKES

{How to not make some}

SCHEMA DESIGN

EYE-CATCHING
GOOD USE OF NEGATIVE SPACE
CONTRAST
ARRANGEMENT
EVOCATIVE OF THE STRUGGLE OF MODERN MAN



LESSONS IN SCHEMA DESIGN

LEARN FROM MY MISTAKES

About me

- * Open source software engineer
- * Postgres fan
- * Over decade of professional DB-driven app development
- * Co-founder and CTO of DoctorBase
- * Lived with my bad decisions

Topics

- * Normalization
- * Dumb schema design
- * Good schema design
- * Data integrity
- * Pro tips

Normalization

- * Eliminating data redundancy
- * Prevent update anomalies
- * Have a unique way to refer to rows for update/delete
- * Great with PostgreSQL foreign keys

DUMB SCHEMA

Table "postgresql_talks"		
Column	Type	Modifiers
talk	text	
speaker	text	
time	timestamp without time zone	default now()

- Dumb table name
- No primary key ID
- Speaker should be foreign key
- Timezone-dependent



KEYS

```
CREATE TABLE postgresql_talk (
    id SERIAL PRIMARY KEY,
```

- Composite PKs - DUMB
- External system PK - DUMB
- No PK - EXTRA DUMB
- Lots of FKs - SLOW DELETES
- Josh Berkus dislikes surrogate keys 😞

DUMB SCHEMA

```
dumb=# insert into postgresql_talks values ('How to suck at RDBMSes', 'Mischa Spiegelmock');

dumb=# insert into postgresql_talks values ('Pg LISTEN/NOTIFY dopeness', 'Misha Spiegelmock');

dumb=# select * from postgresql_talks;
      name       |      speaker      |          time
-----+-----+-----+
 How to suck at RDBMSes | Mischa Spiegelmock | 2016-01-25 23:46:49.659108
 Pg LISTEN/NOTIFY dopeness | Misha Spiegelmock | 2016-01-25 23:47:14.300665
(2 rows)
```

LESS DUMB SCHEMA

```
CREATE TABLE speaker (
    id SERIAL PRIMARY KEY,
    created TIMESTAMP NOT NULL DEFAULT NOW(),
    name TEXT
);
```

```
CREATE TABLE talk (
    id SERIAL PRIMARY KEY,
    created TIMESTAMPTZ NOT NULL DEFAULT NOW(),
    name TEXT,
    speaker INTEGER NOT NULL REFERENCES speaker(id) ON DELETE CASCADE,
    time TIMESTAMP NOT NULL
);
```

```
ALTER TABLE talk ADD CONSTRAINT "table_unique_talk_speaker" UNIQUE
(name, speaker);
```

LESS DUMB SCHEMA

```
dumb=# INSERT INTO speaker (name) VALUES ('Mischa Spiegelmock');
INSERT 0 1
dumb=# INSERT INTO TALK (name, speaker, time) VALUES ('Schema Design', 1, '2016-02-10');
INSERT 0 1
```

```
dumb=# SELECT * FROM TALK;
+-----+-----+-----+-----+-----+
| id | created | name | speaker | time |
+-----+-----+-----+-----+-----+
| 1 | 2016-02-09 18:08:08.592791 | Schema Design | 1 | 2016-02-10 00:00:00 |
+-----+-----+-----+-----+-----+
(1 row)
```

```
dumb=# DELETE FROM speaker WHERE id=1;
DELETE 1
```

```
dumb=# SELECT * FROM TALK;
+-----+-----+-----+-----+-----+
| id | created | name | speaker | time |
+-----+-----+-----+-----+-----+
(0 rows)
```

TALK REVIEWS

```
CREATE TABLE talk_review (
    id SERIAL PRIMARY KEY,
    created TIMESTAMPTZ NOT NULL DEFAULT NOW(),
    talk_id INTEGER REFERENCES talk(id),
    recommendable BOOLEAN,
    email TEXT,
    rating INTEGER
);
ALTER TABLE talk_review ADD CONSTRAINT "review_valid_rating"
    CHECK (rating IS NULL OR rating > 0 AND rating <= 5);
```

MATERIALIZED VIEW REPORTS

```
CREATE MATERIALIZED VIEW review_stat AS
SELECT
    talk_id,
    COUNT( * ) AS total_reviews,
    ROUND( AVG( rating ), 2 ) AS talk_rating,
    SUM( CASE WHEN email IS NOT NULL THEN 1 ELSE 0 END ) AS
total_reviews_with_email,
    ROUND( AVG( recommendable ), 2 ) AS avg_recommendable,
FROM talk_review WHERE rating IS NOT NULL
GROUP BY talk_id ORDER BY talk_id;
```

```
REFRESH MATERIALIZED VIEW review_stat;
```

REFRESH MATERIALIZED VIEW

```
> INSERT INTO talk_review (talk_id, recommendable, email, rating) VALUES (1, 'f', NULL, 3);
> INSERT INTO talk_review (talk_id, recommendable, email, rating) VALUES (1, 't', 'me@foo', 1);
> INSERT INTO talk_review (talk_id, recommendable, email, rating) VALUES (1, 'f', NULL, 4);

> SELECT * FROM review_stat;
 talk_id | total_reviews | talk_rating | total_reviews_with_email | avg_recommendable
-----+-----+-----+-----+
(0 rows)

> REFRESH MATERIALIZED VIEW review_stat;
REFRESH MATERIALIZED VIEW
> SELECT * FROM review_stat;
 talk_id | total_reviews | talk_rating | total_reviews_with_email | avg_recommendable
-----+-----+-----+-----+
 1 | 3 | 2.6666666666666667 | 1 | 0.3333333333333333
(1 row)
```

TIMEZONES

```
dumb=# show timezone;
 TimeZone
-----
 US/Pacific
(1 row)

dumb=# insert into postgresql_talks values ('How to suck at RDBMSes', 'Mischa
Spiegelmock');
dumb=# insert into postgresql_talks values ('Pg LISTEN/NOTIFY dopeness', 'Mischa
Spiegelbokkers');
dumb=# select * from postgresql_talks;
      name       |      speaker      |          time
-----+-----+-----+
 How to suck at RDBMSes | Mischa Spiegelmock | 2016-01-25 23:46:49.659108
 Pg LISTEN/NOTIFY dopeness | Mischa Spiegelbokkers | 2016-01-25 23:47:14.300665
(2 rows)
```

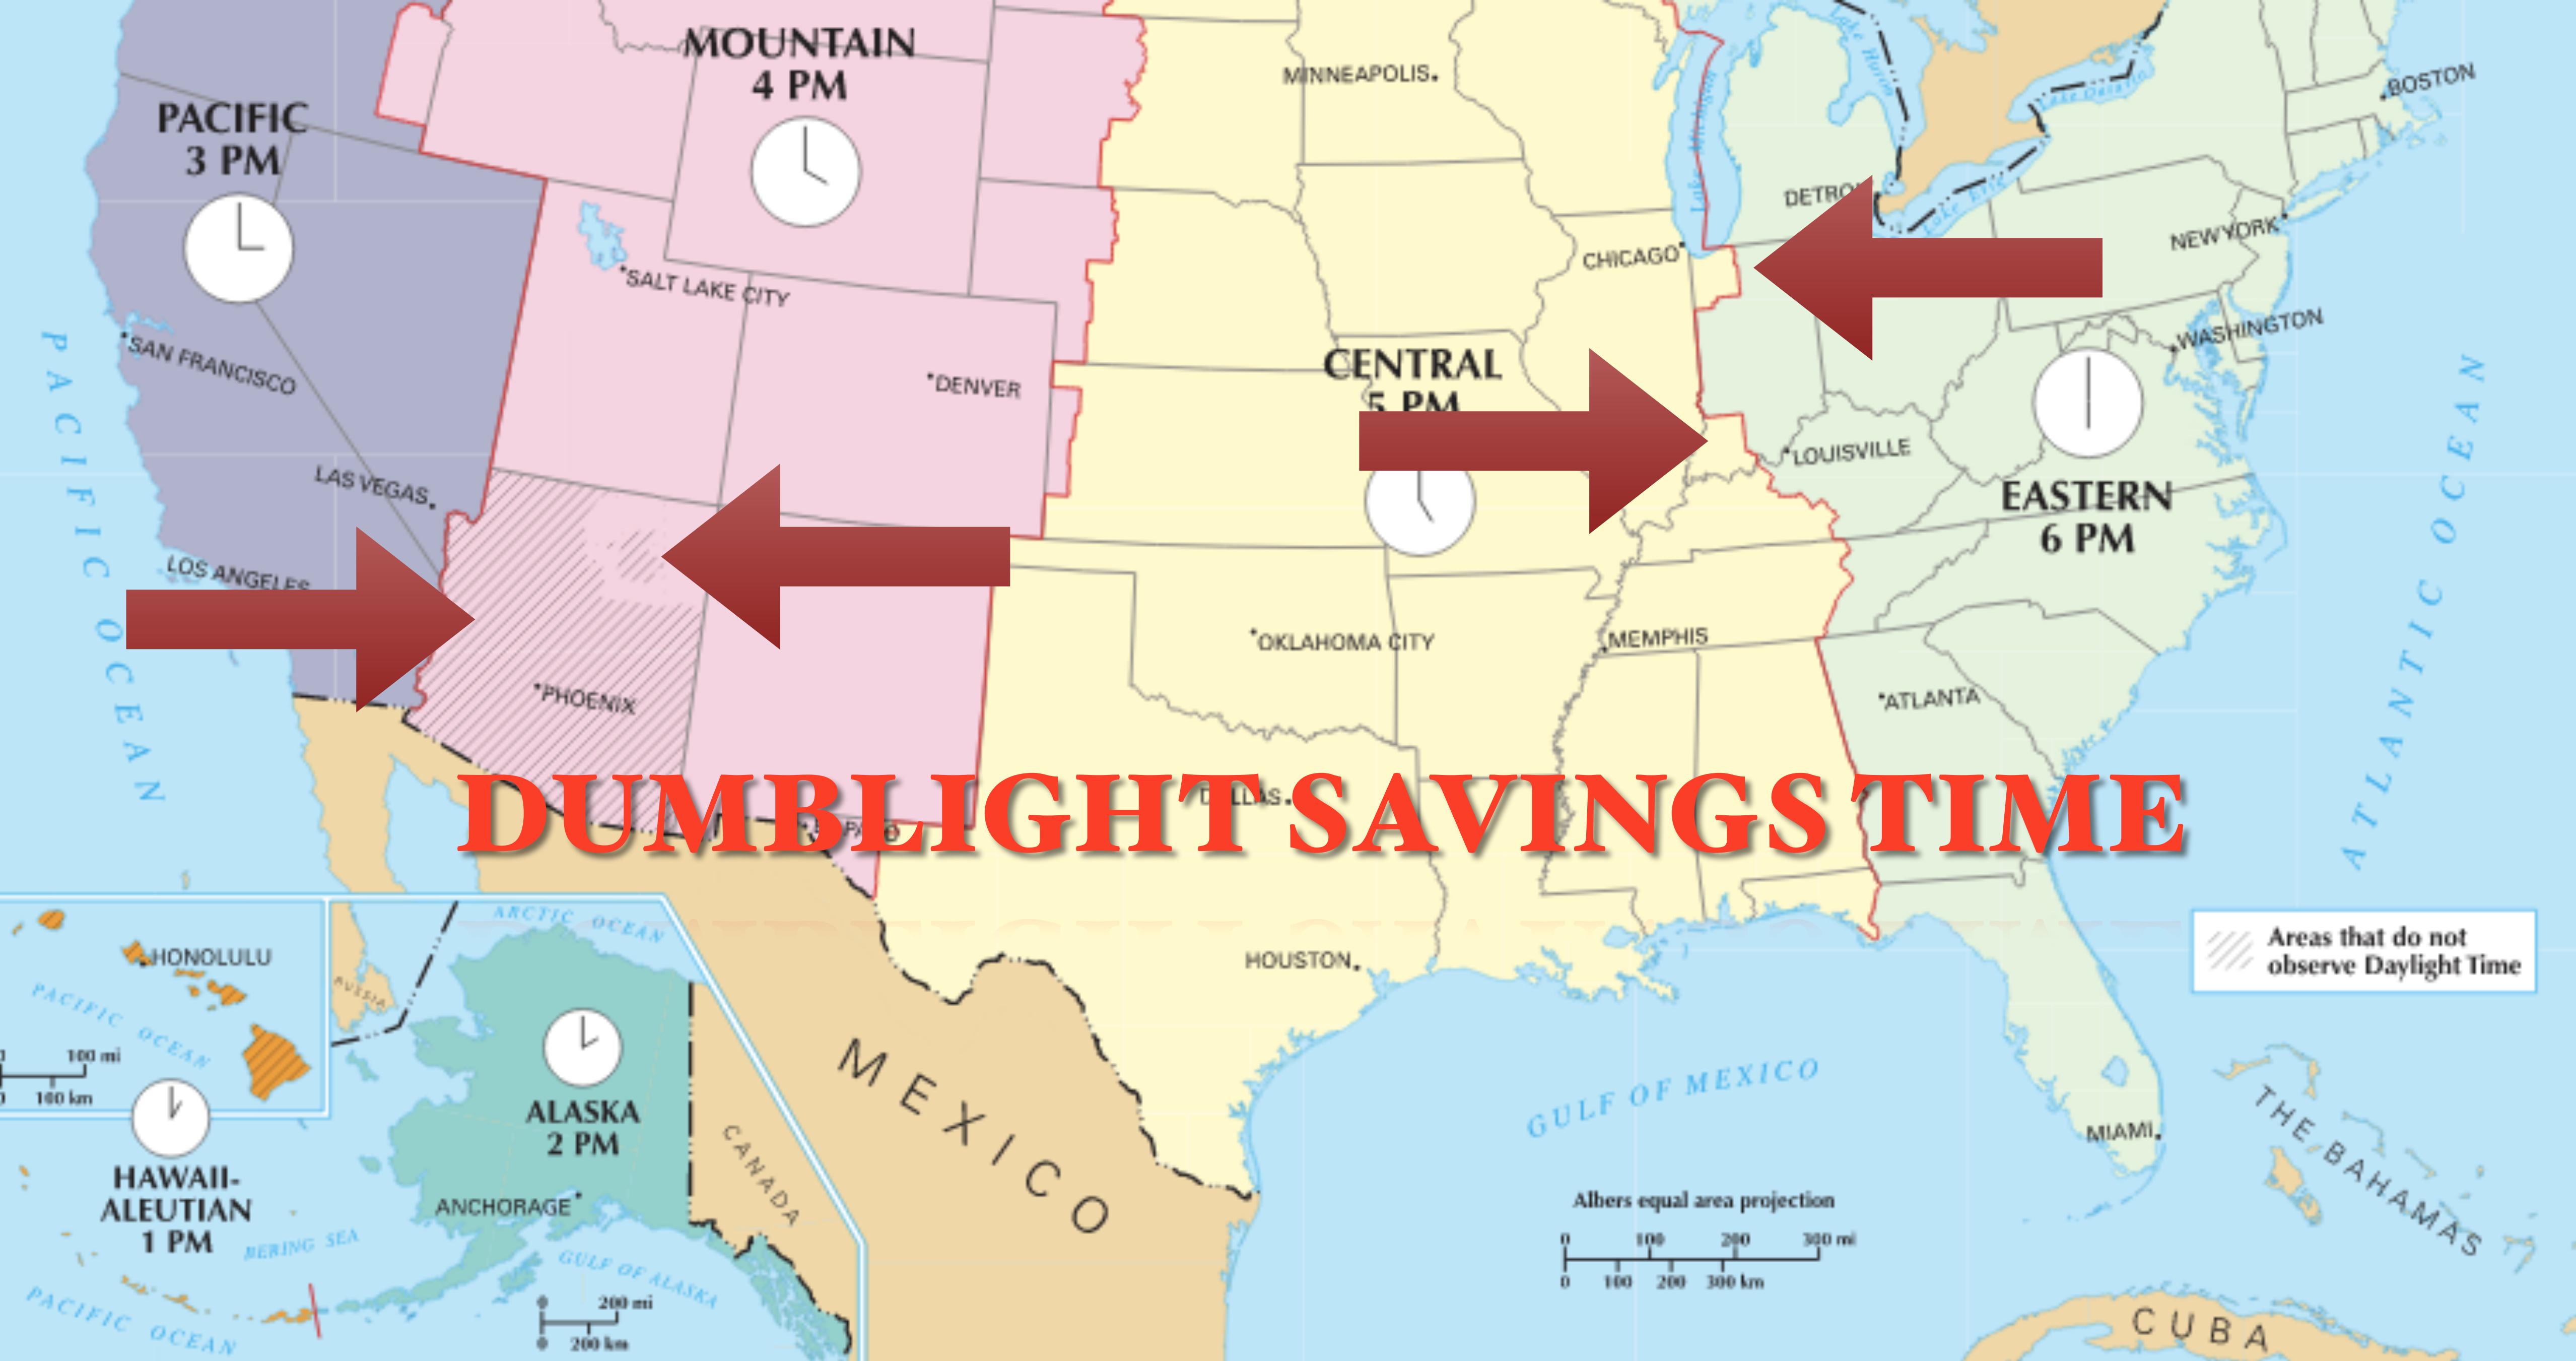
TIMEZONES

```
dumb=# set timezone='US/Hawaii';
dumb=# insert into postgresql_talks values ('Time zones suck', 'Mischa
Spiegelmock');
INSERT 0 1
dumb=# select * from postgresql_talks;
```

name	speaker	time
How to suck at RDBMSes	Mischa Spiegelmock	2016-01-25 23:46:49.659108
Pg LISTEN/NOTIFY dopeness	Mischa Spiegelbokkers	2016-01-25 23:47:14.300665
Time zones suck	Mischa Spiegelmock	2016-01-25 21:48:34.619602
(3 rows)		

TIMEZONE DUMBNESS

- “I’ll just use the system timezone”
- “I can store a timezone with a GMT offset”
- “I can move Pg to another machine and everything will stay the same”
- “I can plant more potatoes with DST”
- “Daylight savings time is regular”



#	Rule	NAME	FROM	TO	TYPE	IN	ON	AT	SAVE	LETTER/S
Rule	US	1918	1919	-		Mar	lastSun	2:00	1:00	D
Rule	US	1918	1919	-		Oct	lastSun	2:00	0	S
Rule	US	1942	only	-		Feb	9	2:00	1:00	W # War
Rule	US	1945	only	-		Aug	14	23:00u	1:00	P # Peace
Rule	US	1945	only	-		Sep	30	2:00	0	S
Rule	US	1967	2006	-		Oct	lastSun	2:00	0	S
Rule	US	1967	1973	-		Apr	lastSun	2:00	1:00	D
Rule	US	1974	only	-		Jan	6	2:00	1:00	D
Rule	US	1975	only	-		Feb	23	2:00	1:00	D
Rule	US	1976	1986	-		Apr	lastSun	2:00	1:00	D
Rule	US	1987	2006	-		Apr	Sun>=1	2:00	1:00	D
Rule	US	2007	max	-		Mar	Sun>=8	2:00	1:00	D
Rule	US	2007	max	-		Nov	Sun>=1	2:00	0	S
....										
#	Rule	NAME	FROM	TO	TYPE	IN	ON	AT	SAVE	LETTER
Rule	NYC	1920	only	-		Mar	lastSun	2:00	1:00	D
Rule	NYC	1920	only	-		Oct	lastSun	2:00	0	S
Rule	NYC	1921	1966	-		Apr	lastSun	2:00	1:00	D
Rule	NYC	1921	1954	-		Sep	lastSun	2:00	0	S
Rule	NYC	1955	1966	-		Oct	lastSun	2:00	0	S
#	Zone	NAME	GMTOFF	RULES	FORMAT	[UNTIL]				
Zone	America/New_York		-4:56:02	-	LMT	1883 Nov 18	12:03:58			
			-5:00	US	E&ST	1920				
			-5:00	NYC	E&ST	1942				
			-5:00	US	E&ST	1946				
			-5:00	NYC	E&ST	1967				
			-5:00	US	E&ST					

Maintainer: Paul Eggert

<END TIMEZONE RANT>

Sane Dates and Times

- * Store everything in GMT or Unix Timestamp
- * Use type: `TIMESTAMPTZ`
- * Convert to local time on display
- * Store TZ as IANA name (“America/Los_Angeles”)
- * E-Z date math (`1day = 86400`, `1hr = 3600`, etc)
- * Never store age, store birthdate

.CREATED

```
CREATE TABLE postgresql_talk (
    id SERIAL PRIMARY KEY,
    created TIMESTAMPTZ NOT NULL DEFAULT NOW()
```

- **Reporting**
- **Fix blunders in a date range**

CONSTRAINTS

```
ALTER TABLE foo ADD CONSTRAINT "must_choose_one" CHECK ((A is null)::int + (B is null)::int + (C is null)::int = 2);
```

```
CREATE TYPE foo_opt AS ENUM ('A', 'B', 'C');
```

```
CREATE TYPE sex AS ENUM ('male', 'female', 'other');
```

TRIGGER WARNING: TRIGGERS

TSVECTOR TRIGGERS

```
CREATE FUNCTION user_keyword_update() RETURNS trigger
    LANGUAGE plpgsql
AS $$

BEGIN
    NEW.keywords := to_tsvector('pg_catalog.english',
        NEW.id || ' ' || coalesce(NEW.username, '') || ' ' ||
        coalesce(NEW.email_address, '') || ' ' ||
        coalesce(NEW.name, '') || ' ' ||
        coalesce(NEW.phone, '') || ' '
    );
    RETURN NEW;
END;

$$;
```

```
CREATE TRIGGER user_keyword_update_trigger BEFORE INSERT OR UPDATE
ON "user" FOR EACH ROW EXECUTE PROCEDURE user_keyword_update();
```

LOCATION UPDATE TRIGGERS

```
CREATE OR REPLACE FUNCTION update_georeferenced_table() RETURNS
TRIGGER AS $$

DECLARE
    name TEXT;
BEGIN
    name := TG_ARGV[0];
    IF ( (TG_OP = 'INSERT' AND NEW.location IS NOT NULL) OR (TG_OP =
    'UPDATE' AND NEW.location IS DISTINCT FROM OLD.location) ) THEN
        NEW.updated_location := current_timestamp;
    END IF;
    PERFORM pg_notify(name || '_updated', '{
        "id": ' ||
        CAST(NEW.id AS TEXT) ||
        ', "location": ' ||
        ST_AsGeoJSON(NEW.location) ||
        '}'
    );
    RETURN NEW;
END;
```

LISTEN/NOTIFY  WebSocket

github.com/doctorbase/wsnotify

github.com/revmischa/pgnotify-demos

ULTIMATE KILLER SCHEMA

“I’m going to make the `ultimo` super flexible schema and let my users choose what fields to store”

```
CREATE TABLE user_field (
    user_id integer NOT NULL,
    field_id integer NOT NULL,
    value text
);

CREATE TABLE field (
    id integer SERIAL PRIMARY KEY,
    name text
);
```



EAV

Key-value pair store: JSONB

Join the No-noSQL movement

Example	Example Result
<code>to_json('Fred said "Hi."'::text)</code>	<code>"Fred said \"Hi.\\\""</code>
<code>array_to_json('{{1,5}, {99,100}}'::int[])</code>	<code>[[1,5], [99,100]]</code>
<code>row_to_json(row(1,'foo'))</code>	<code>{"f1":1,"f2":"foo"}</code>
<code>json_build_array(1,2,'3',4,5)</code>	<code>[1, 2, "3", 4, 5]</code>
<code>json_build_object('foo',1,'bar',2)</code>	<code>{"foo": 1, "bar": 2}</code>
<code>json_object('{a, 1, b, "def", c, 3.5}')</code>	
<code>json_object('{{a, 1},{b, "def"}, {c, 3.5}}')</code>	<code>{"a": "1", "b": "def", "c": "3.5"}</code>
<code>json_object('{a, b}', '{1,2}')</code>	<code>{"a": "1", "b": "2"}</code>

Object Relational Mappers

- * Define classes for tables
- * Instance == row
- * Hibernate, DBIx::Class, Django-ORM, SQLAlchemy

```
SELECT * FROM talk t LEFT OUTER JOIN talk_review tr ON  
tr.talk_id=t.id WHERE t.id=1;
```

```
my @reviews = $talk->reviews;
```

ORM

```
select team_id from user_team  
where user_id in ($self.id, $user.id)  
group by team_id  
having count(*) > 1;
```

```
query = db.session  
    .query(func.count(UserTeam.team_id))  
    .filter(UserTeam.user_id.in_([ self.id, user.id ]))  
    .group_by(UserTeam.team_id)  
    .having(func.count(UserTeam.user_id) > 1)
```



ORM - POSTGIS

```
class PointGeography(types.UserDefinedType):
    def get_col_spec(self):
        return "GEOMETRY"

    def column_expression(self, col):
        return ST_AsGeoJSON(col, type_=self)

class GeoReferenced():
    updated_location = Column(DateTime())
    location_accuracy_meters = Column(Numeric())
    location = Column(PointGeography)

    def set_location(self, lng, lat):
        if lat is None or lng is None:
            self.location = None
        else:
            self.location = "POINT(%0.16f %0.16f)" %
(float(lng), float(lat))
        db.session.commit()

@property
def lat(self):
    if self.location is None:
        return None
    return parse_point(self.location)[1]

@property
def lng(self):
    if self.location is None:
        return None
    return parse_point(self.location)[0]
```

```
def parse_point(point):
    geo = json.loads(point)
    return geo['coordinates']

@classmethod
def within_clause(cls, latitude, longitude, distance):
    attr = '%s.location' % cls.__tablename__
    point = 'POINT(%0.8f %0.8f)' % (longitude, latitude)
    location = "ST_GeographyFromText(E'SRID=4326;%s')" % point
    return 'ST_DWithin(%s, %s, %d)' % (attr, location, distance)
```

TIMEOUT

```
SET statement_timeout = 120000
```

BLOBS / BYTEA

Never worked well for me.

VARCHAR(255)

Use TEXT
VARCHAR(NN) is pointless

NAMING

No plurals!

USER TABLE

Don't call it “user”

COMPARTMENTALIZATION

Make schemas for related data

MODEL DATA ON THE REAL WORLD

Schema should map real things

WHEN TO USE STORED PROCEDURES?

Thoughts?

Storing Static Files

- * Amazon S3
- * Security tokens
- * Static web hosting
- * Fancy ACLs
- * Requires application-level support
- * Store s3keyname column (and MIME, size, views, s3bucket, ...)

HELP!

irc.freenode.net #postgresql

Slack: <https://denish1.typeform.com/to/S7h9jl>